



Institute for Advanced Studies  
in Basic Sciences  
Gava Zang, Zanjan, Iran

# Introduction to Machine Learning

By Dr. Parvin Razzaghi ([p.razzaghi@iasbs.ac.ir](mailto:p.razzaghi@iasbs.ac.ir))

Winter 2024





# Slide credit

- Professor
  - Dr. Parvin Razzaghi
- Website
  - <http://www.seas.upenn.edu/~cis519>
- Original Slides:
  - Eric Eaton
- Acknowledgment
  - These slides have been prepared based on the materials created by Eric Eaton, with sincere appreciation for the many educators who have generously shared their course materials online. These slides may be reused or adapted for academic purposes, provided that appropriate attribution is given. For any comments or suggestions, please contact: [p.razzaghi@iasbs.ac.ir](mailto:p.razzaghi@iasbs.ac.ir)



# Course Information

- Slides and Assignments:
  - Accessible through the group:



- Quera:
  - [https://quera.org/course/add\\_to\\_course/course/20341/](https://quera.org/course/add_to_course/course/20341/)
  - Password: ml1403



# Communication

- **Contacting the Professor:**

- [p.razzaghi@iasbs.ac.ir](mailto:p.razzaghi@iasbs.ac.ir)


- **Contacting the Teaching Assistants (TAs):**

- TAs' Information:


- Elaheh Sheikh: [[elahehsheikhi@iasbs.ac.ir](mailto:elahehsheikhi@iasbs.ac.ir)] | [@elisheikhi]
- Maryam Lotfi: [[maryam.lotfi@iasbs.ac.ir](mailto:maryam.lotfi@iasbs.ac.ir)] | [@marymltf]
- Reza Pishkar: [[r.pishkar@iasbs.ac.ir](mailto:r.pishkar@iasbs.ac.ir)] | [@rpishkr]
- Mohammadreza Naderlou: [[m.naderlo@iasbs.ac.ir](mailto:m.naderlo@iasbs.ac.ir)] | [@mimre\_non]



## Traditional Programming



## Machine Learning





# When Do We Use Machine Learning?


- ML is used when:
  - Human expertise does not exist (navigating on Mars)
  - Humans can't explain their expertise (speech recognition)
  - Models must be customized (personalized medicine)
  - Models are based on huge amounts of data (genomics)





A classic example of a task that requires machine learning:

It is very hard to say what makes a 2





## Some more examples of tasks that are best solved by using a learning algorithm

- Recognizing patterns:
  - Facial identities or facial expressions
  - Handwritten or spoken words
  - Medical images
- Generating patterns:
  - Generating images or motion sequences
- Recognizing anomalies:
  - Unusual credit card transactions
  - Unusual patterns of sensor readings in a nuclear power plant
- Prediction:
  - Future stock prices or currency exchange rates





# Sample Applications

- Web search
- Computational biology
- Finance
- E-commerce
- Space exploration
- Robotics
- Information extraction
- Social networks
- Debugging software
- ....



# Samuel's Checkers-Player

“Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.”

Arthur  
Samuel  
(1959)





# What is Machine Learning?

“Learning is any process by which a system improves performance from experience.”

Herbert  
Simon

A red speech bubble containing the name 'Herbert Simon'.

- Definition by Tom Mitchell (1998):
  - Machine Learning is the study of algorithms that
    - improve their performance  $P$
    - at some task  $T$
    - with experience  $E$ .
  - A well-defined learning task is given by  $\langle P, T, E \rangle$ .



# Defining the Learning Task

Improve on task  $T$ , with respect to performance metric  $P$ , based on experience  $E$

|                                                                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>• <math>T</math>: Playing checkers</li><li>• <math>P</math>: Percentage of games won against an arbitrary opponent</li><li>• <math>E</math>: Playing practice games against itself</li></ul>                                                                                 | <ul style="list-style-type: none"><li>• <math>T</math>: Recognizing hand-written words</li><li>• <math>P</math>: Percentage of words correctly classified</li><li>• <math>E</math>: Database of human-labeled images of handwritten words</li></ul>                        |
| <ul style="list-style-type: none"><li>• <math>T</math>: Driving on four-lane highways using vision sensors</li><li>• <math>P</math>: Average distance traveled before a human-judged error</li><li>• <math>E</math>: A sequence of images and steering commands recorded while observing a human driver.</li></ul> | <ul style="list-style-type: none"><li>• <math>T</math>: Categorize email messages as spam or legitimate.</li><li>• <math>P</math>: Percentage of email messages correctly classified.</li><li>• <math>E</math>: Database of emails, some with human-given labels</li></ul> |




# State of the Art Applications of Machine Learning



# Autonomous Cars

- Nevada made it legal for autonomous cars to drive on roads in June 2011
- As of 2013, four states (Nevada, Florida, California, and Michigan) have legalized autonomous cars.


Penn's Autonomous  
Car (Ben Franklin  
Racing Team)





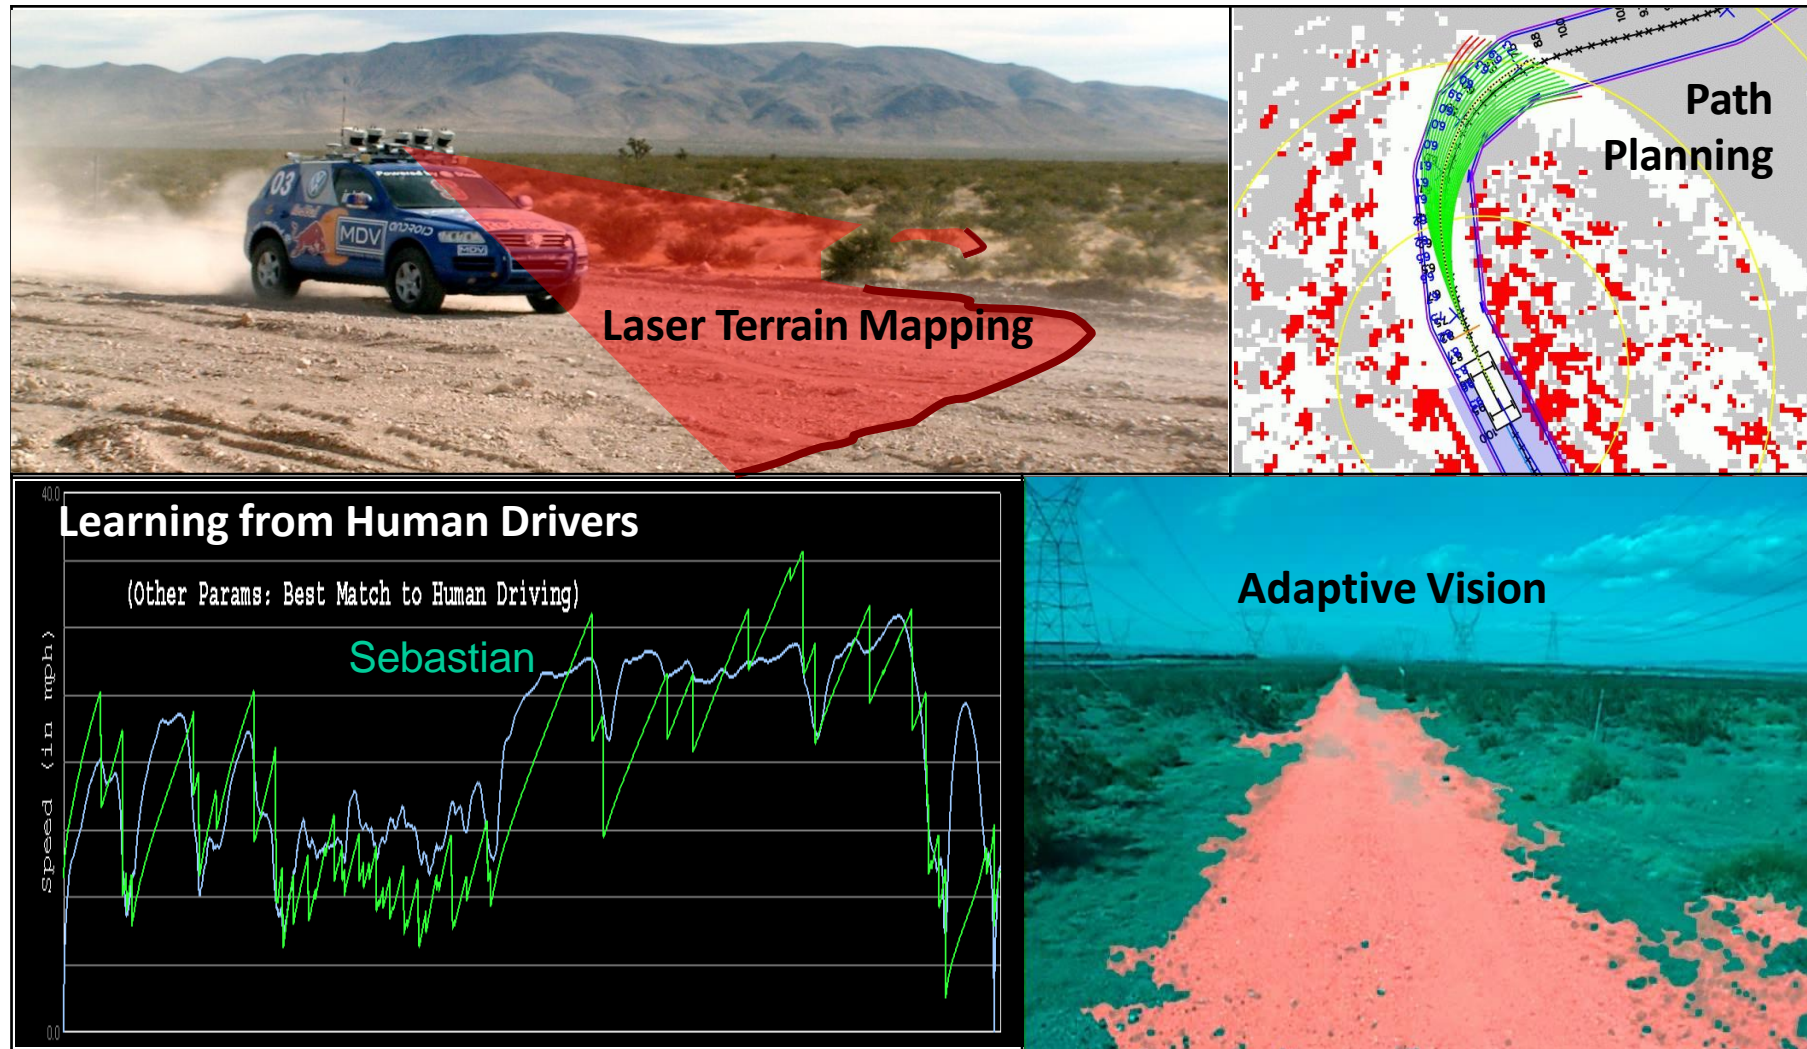


# Autonomous Car Sensors





# Autonomous Car Technology







# Deep Learning in the Headlines


BUSINESS NEWS

**MIT Technology Review**

## Is Google Cornering the Market on Deep Learning?

A cutting-edge corner of science is being wooed by Silicon Valley, to the dismay of some academics.

By Antonio Regalado on January 29, 2014



How much are a dozen deep-learning researchers worth? Apparently, more than \$400 million.

This week, Google reportedly paid that much to acquire [DeepMind Technologies](#), a startup based in London.

This is Freescale *make it*

**BloombergBusinessweek Technology**

Acquisitions

## The Race to Buy the Human Brains Behind Deep Learning Machines

By Ariane Vance | January 27, 2014

intelligence projects. "DeepMind is bona fide in terms of its research capabilities and depth," says Peter Lee, who heads Microsoft Research.

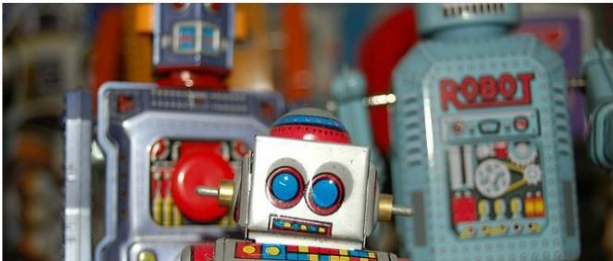
According to Lee, Microsoft, Facebook (FB), and Google find themselves in a battle for deep learning talent. Microsoft has gone from four full-time deep learning experts to 70 in the past three years. "We would have more if the talent was there to


**WIRED** GEAR SCIENCE ENTERTAINMENT BUSINESS SECURITY DESIGN

INNOVATION INSIGHTS community content featured

## Deep Learning's Role in the Age of Robots

BY JULIAN GREEN, JETPAC 05.02.14 2:56 PM






**DEEP LEARNING**

- » Computers learning and growing on their own
- » Able to understand complex, massive amounts of data

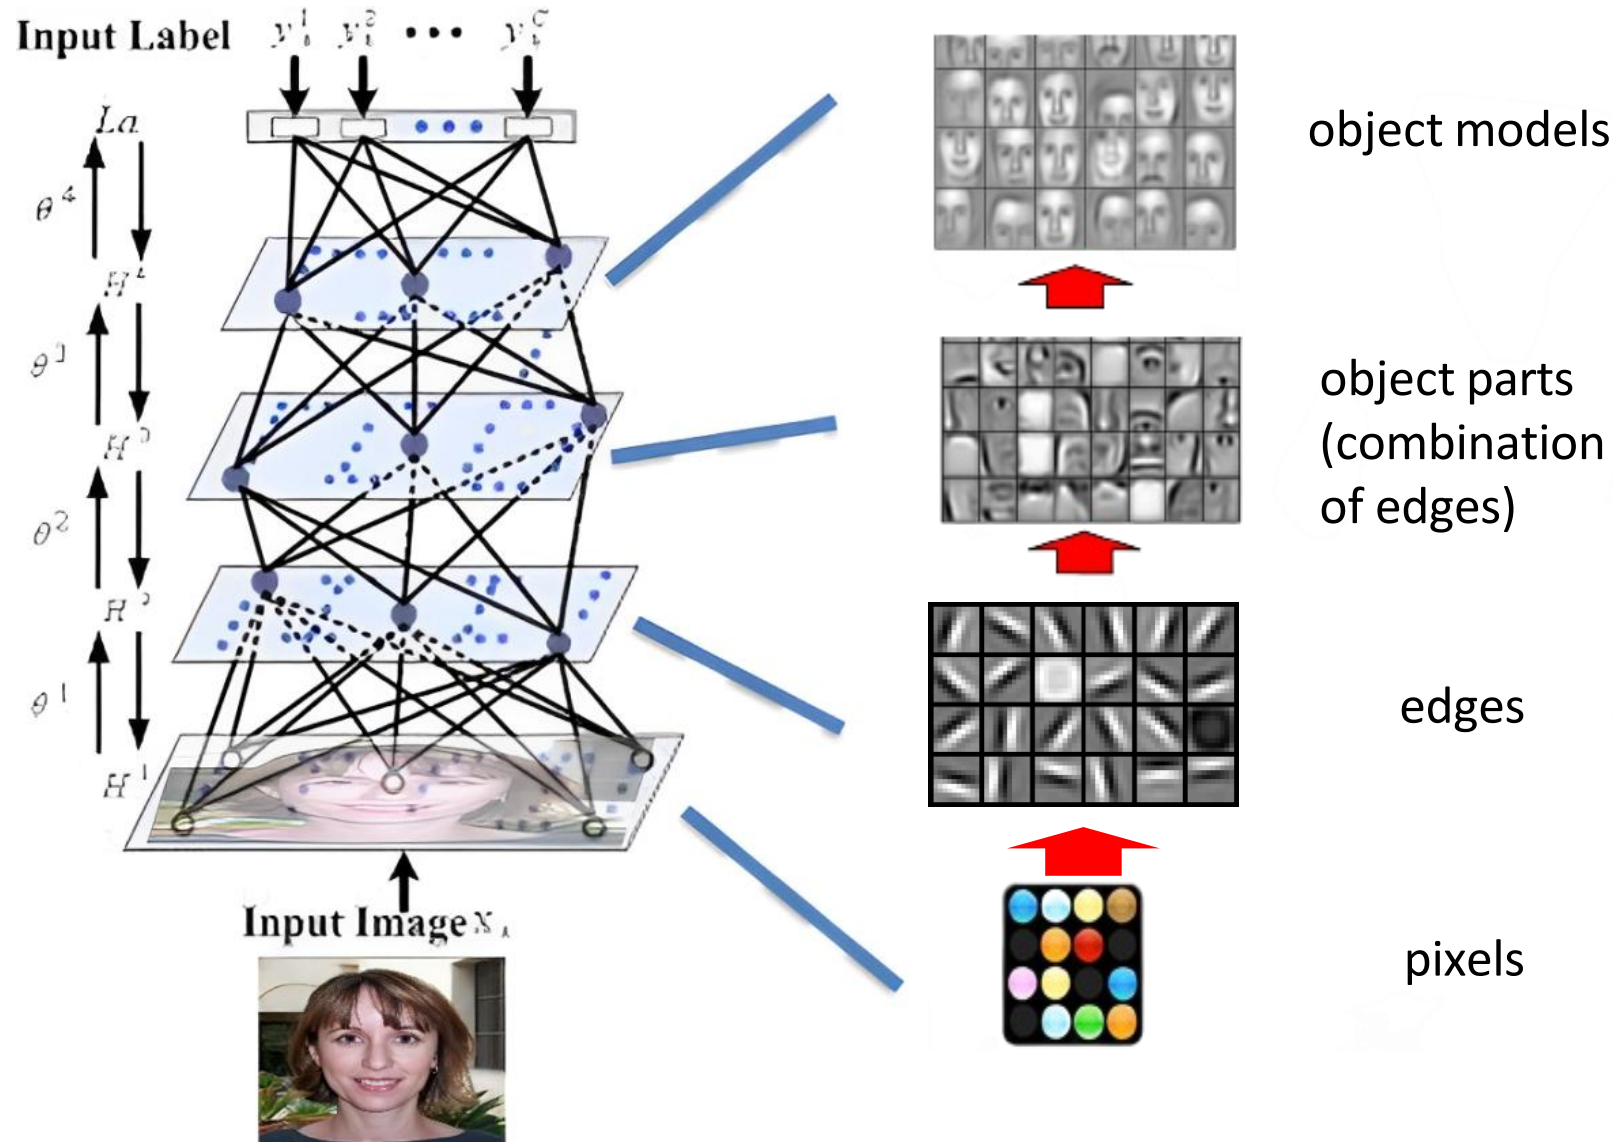
DATA ECONOMY  
DEEP LEARNING

BROUGHT TO YOU BY: 

**CNBC**







# Deep Belief Net on Face Images





# Training on Multiple Objects

- Trained on 4 classes (cars, faces, motorbikes, airplanes).
- Second layer: Shared-features and object-specific features.
- Third layer: More specific features.






# Scene Labeling via Deep Learning






# Machine Learning in Automatic Speech Recognition

## A Typical Speech Recognition System



## ML used to predict of phone states from the sound spectrogram



Deep learning has state-of-the-art results


| # Hidden Layers   | 1    | 2    | 4    | 8    | 10   | 12   |
|-------------------|------|------|------|------|------|------|
| Word Error Rate % | 16.0 | 12.8 | 11.4 | 10.9 | 11.0 | 11.1 |

Baseline GMM performance = 15.4%

[Zeiler et al. "On rectified linear units for speech recognition" ICASSP 2013]



# Impact of Deep Learning in Speech Technology





# Types of Learning





# Types of Learning


- **Supervised (inductive) learning**
  - Given: training data + desired outputs (labels)
- **Unsupervised learning**
  - Given: training data (without desired outputs)
- **Semi-supervised learning**
  - Given: training data + a few desired outputs
- **Reinforcement learning**
  - Rewards from sequence of actions





# Supervised Learning: Regression


- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is real-valued == regression





# Supervised Learning: Classification


- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is categorical == classification





# Supervised Learning: Classification


- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is categorical == classification





# Supervised Learning

- $x$  can be multi-dimensional
  - Each dimension corresponds to an attribute




- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...



# Unsupervised Learning


- Given  $x_1, x_2, \dots, x_n$  (without labels)
- Output hidden structure behind the  $x$ 's
  - E.g., clustering





# Unsupervised Learning

Genomics application: group individuals by genetic similarity






# Unsupervised Learning



Organize computing clusters



Social network analysis



Market segmentation



Astronomical data analysis




# Reinforcement Learning

- Given a sequence of states and actions with (delayed) rewards, output a policy
  - Policy is a mapping from states  $\rightarrow$  actions that tells you what to do in a given state
- Examples:
  - Credit assignment problem
  - Game playing
  - Robot in a maze
  - Balance a pole on your hand





# The Agent-Environment Interface




Agent and environment interact at discrete time steps :  $t = 0, 1, 2, K$

Agent observes state at step  $t$ :  $s_t \in S$

produces action at step  $t$ :  $a_t \in A(s_t)$

gets resulting reward :  $r_{t+1} \in \mathcal{R}$

and resulting next state :  $s_{t+1}$





# Reinforcement Learning



<https://www.youtube.com/watch?v=4cgWya-wjgY>

# Inverse Reinforcement Learning

- Learn policy from user demonstrations



Stanford Autonomous Helicopter

<http://heli.stanford.edu/>

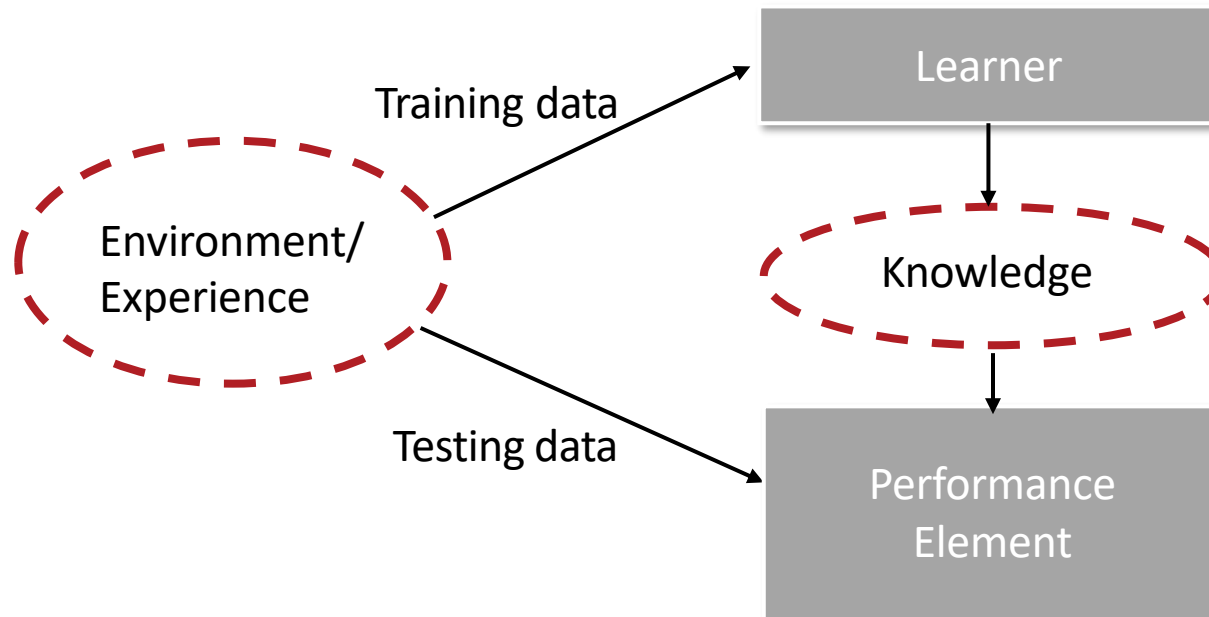
<https://www.youtube.com/watch?v=VCdxqn0fcnE>



# Framing a Learning Problem

# Designing a Learning System

- Choose the training experience
- Choose exactly what is to be learned
  - i.e. the **target function**
- Choose how to represent the target function
- Choose a learning algorithm to infer the target function from the experience





# Training vs. Test Distribution

- We generally assume that the training and test examples are independently drawn from the same overall distribution of data
  - We call this “i.i.d” which stands for “independent and identically distributed”
- If examples are not independent, requires *collective classification*
- If test distribution is different, requires *transfer learning*



# ML in a Nutshell

- Tens of thousands of machine learning algorithms
  - Hundreds new every year
- Every ML algorithm has three components:
  - **Representation**
  - **Optimization**
  - **Evaluation**





# Various Function Representations

- Numerical functions
  - Linear regression
  - Neural networks
  - Support vector machines
- Symbolic functions
  - Decision trees
  - Rules in propositional logic
  - Rules in first-order predicate logic
- Instance-based functions
  - Nearest-neighbor
  - Case-based
- Probabilistic Graphical Models
  - Naïve Bayes
  - Bayesian networks
  - Hidden-Markov Models (HMMs)
  - Probabilistic Context Free Grammars (PCFGs)
  - Markov networks





# Various Search/Optimization Algorithms

- Gradient descent
  - Perceptron
  - Backpropagation
- Dynamic Programming
  - HMM Learning
  - PCFG Learning
- Divide and Conquer
  - Decision tree induction
  - Rule learning
- Evolutionary Computation
  - Genetic Algorithms (GAs)
  - Genetic Programming (GP)
  - Neuro-evolution




# Evaluation

- Accuracy
- Precision and recall
- Squared error
- Likelihood
- Posterior probability
- Cost / Utility
- Margin
- Entropy
- K-L divergence
- etc.



# ML in Practice



- Understand domain, prior knowledge, and goals
- Data integration, selection, cleaning, pre-processing, etc.
- Learn models
- Interpret results
- Consolidate and deploy discovered knowledge



# Lessons Learned about Learning

- Learning can be viewed as using direct or indirect experience to approximate a chosen target function.
- Function approximation can be viewed as a search through a space of hypotheses (representations of functions) for one that best fits a set of training data.
- Different learning methods assume different hypothesis spaces (representation languages) and/or employ different search techniques.

**Any** questions?

